

Project Submission 8
Video Classification

1. Topic

The goal of this project is to build a neural network that is capable of detecting human doing housework. Housework activities include cleaning, cooking, washing dishes, etc. For this project, I narrow down the definition of housework to two activities: **house cleaning and washing dishes**.

2. What's new in this submission

2.1. The model

Surprisingly, the modified neural network model in submission 6 performed poorly on untrained data, even though it got a very good testing accuracy:

- Accuracy when tested with the sample video provided on YouTube by the professor: **0%**.
- Accuracy when tested with 19 unmet/untrained activities: **50%**.

Thus, for this final submission, I made several changes in the neural network architecture and the dataset. Three most significant modifications made on the neural network architecture are:

- The number of layers and neurons per layer are changed.
- Increase the number of epochs to 25.
- Changed the loss function to *categorical* from *binary* cross-entropy.

Input	Hidden	Hidden	Hidden	Hidden	Hidden	Output
25088	1024	512	256	256	128	3

The training dataset has been increased to include 5 more indoor activities. The architecture of the neural network has also been improved to be able to learn from this more complex dataset. I've also found that the best number of epochs increase from 15 to 25. This gives more time for the model to learn.

In last submission, the output of the neural network is binary: housework or not_housework. This setup is problematic because we have two distinct housework activities: house cleaning and washing dishes. Combining them into a group makes it hard for the neural network to detect because it is difficult to find the common between the two activities.

Thus, in this submission, the output is changed to *categorical*: house cleaning, washing dishes and not housework.

With this configuration, it is easier for the neural model to learn and therefore increases the prediction accuracy:

- Accuracy when tested with the sample video provided on YouTube by the professor: **100%**.
- Accuracy when tested with 19 unmet/untrained activities: **81%**.

2.2. The training/testing sample videos naming format

In previous submissions, names of the input videos have to be put in *testlist.txt* and *trainlist.txt* in a very specific and complex way (see section 9 in previous submissions). This was because the code was written to directly handle the videos from the UCF101 dataset.

For this submission, the code has been rewritten to handle the full name of input videos. Tester just needs to type in the original name of the video in *testlist.txt* or *trainlist.txt*. No reformatting is needed (see section 9 for more information).

3. Dataset

A total of 44 selected activities in the UCF101 and STAIR datasets are used. The final dataset includes a total of **3,091** videos grouped into 3 sets: *training set*, *testing set*, and *unmet testing set*. The *unmet testing set* is used to test the performance of the model against several activities that it has never encountered or been trained on.

- Training set: **1,622** videos
 - 312 videos of house cleaning
 - 350 videos of washing dishes
 - 960 videos of other 23 indoor activities:
 - Apply Eye Makeup
 - Apply Lipstick
 - Baby Crawling
 - Blow Dry Hair
 - Blow Candles
 - Brushing Teeth
 - Cutting in Kitchen
 - Folding Laundry
 - Ironing
 - Knitting
 - Playing Cello
 - Playing Guitar
 - Playing Piano
 - Playing Violin
 - Pouring Coffee or Tea
 - Pull Ups
 - Push Ups

- Shaving Beard
- Studying
- Telephoning
- Washing Hands
- Wiping Window
- Wearing Shoes
- Training videos are split into **11,166** total frames
 - Training frames: **8,932**
 - Validation frames: **2,234**
- Testing set: **967** videos
- Unmet testing set: **502** videos of 19 untrained activities.
 - Eating meal
 - Exercising
 - Haircutting
 - Lying on floor
 - Punching
 - Reading newspaper
 - Rock climbing indoor
 - Rowing
 - Salsa Spin
 - Smoking
 - Soccer Juggling
 - Taichi
 - Throwing Trash
 - Using Computer
 - Walking with Dog
 - Wall Pushup
 - Washing Face
 - Writing on Board
 - Yoyo

4. DNN Model

4.1. Architecture

Input vector	25088
Hidden layer	1024
Hidden layer	512
Hidden layer	256
Hidden layer	256
Hidden layer	128
Output layer	3

4.2. Input tensor

Initial training: (8932, 224, 224, 3)

Initial validating: (2234, 224, 224, 3)

Initial testing: (7110, 224, 224, 3)

Training and validating data are preprocessed with the pretrained VGG-16 model and then reshaped [1]. Thus, the final training and validating data is:

Final training: (8932, 25088)

Final validating: (2234, 25088)

Final testing: (7110, 25088)

4.3. Output tensor

Output training: (8932, 3)

Output validating: (2234, 3)

Output testing: (7110, 3)

5. **Hyperparameters**

5.1. Range of Hyperparameters Tried

Batch size	64, 128, 256
Epochs	15, 25, 35, 45
Dropout	0.1 - 0.5

5.2. Optimal Hyperparameters:

Batch size	128
Epochs	25
Dropout	0.5

6. **Annotated Code**

- Part of the code is referenced from this article [1].
- Part of them is added/modified to fit the project requirements and optimize the model performance, including:
 - Model architecture.
 - Video frames labeling
 - Output time/label data to JSON file.
 - Testing method (per frame instead of per video as in the reference code)

```

training.py × testing.py
home > phatnguyen > Documents > CSCE-636-Video-Classification > training.py

43 # create a dataframe
44 train = pd.DataFrame()
45 train['video_name'] = videos
46 train = train[:-1]
47 train.head()
48
49 # ===== #
50
51 # remove old frames in the extracted_frames folder
52 files = glob('training_videos/extracted_frames/*.jpg')
53 for f in files:
54     os.remove(f)
55
56 # extract the frames from training videos
57 for i in tqdm(range(train.shape[0])):
58     count = 0
59     videoFile = train['video_name'][i]
60     cap = cv2.VideoCapture('training_videos/' + videoFile) # capturing the video from the given path
61     frameRate = cap.get(5) #frame rate
62     while(cap.isOpened()):
63         frameId = cap.get(1) #current frame number
64         ret, frame = cap.read()
65         if (ret != True):
66             break
67         if (frameId % math.floor(frameRate) == 0): # get one frame per second
68             # storing the frames in a new folder named extracted_frames
69             filename = 'training_videos/extracted_frames/' + videoFile + "_frame%d.jpg" % count; count+=1
70             cv2.imwrite(filename, frame)
71     cap.release()
72
73 # ===== #
74
75 # get label for all images
76 images = glob("training_videos/extracted_frames/*.jpg")
77 train_image = []
78 train_class = []
79 for i in tqdm(range(len(images))):
80     # creating the image name
81     train_image.append(images[i].split('/')[2])
82
83     # creating the class of image
84     if images[i].find('MoppingFloor') != -1:
85         train_class.append("housecleaning")
86     elif images[i].find('WashingDishes') != -1:
87         train_class.append("washingdishes")
88     else:
89         train_class.append("not_housework")
90
91 # storing the images and their class in a dataframe
92 train_data = pd.DataFrame()
93 train_data['image'] = train_image
94 train_data['class'] = train_class
95
96 # convert the dataframe into csv file
97 train_data.to_csv('training_frames_list.csv', header=True, index=False)
98
99 # ===== #

```

Listing 1: Extract the frames from the training dataset and label each frame.

```

130
131 # split the videos into training and validation set
132 y = train['class']
133 x_train, x_validate, y_train, y_validate = train_test_split(x, y, random_state=42, test_size=0.2, stratify = y)
134
135 # create dummies of target variable for train and validation set
136 y_train = pd.get_dummies(y_train)
137 y_validate = pd.get_dummies(y_validate)
138
139 print "y_train shape: ",
140 print(y_train.shape)
141 print "y_validate shape: ",
142 print(y_validate.shape)
143

```

Listing 2: Split all the frames into training set and validation set

```

158 # ===== #
159
160 print "Processing training data through VGG16 ...",
161
162 # create the base model of pre-trained VGG16 model
163 base_model = VGG16(weights='imagenet', include_top=False)
164
165 # extract features for training frames
166 x_train = base_model.predict(x_train)
167
168 # extract features for validation frames
169 x_validate = base_model.predict(x_validate)
170
171 print("Done")
172 print "x_train shape: ",
173 print(x_train.shape)
174 print "x_validate shape: ",
175 print(x_validate.shape)
176
177 # ===== #
178
179 print "Reshaping training data for the final fully connected neural network ... ",
180
181 # reshape the training as well as validation frames in single dimension
182 x_train = x_train.reshape(x_train.shape[0], 7*7*512)
183 x_validate = x_validate.reshape(x_validate.shape[0], 7*7*512)
184
185 # normalize the pixel values
186 max = x_train.max()
187 x_train = x_train/max
188 x_validate = x_validate/max
189
190 print("Done")
191 print "x_train shape: ",
192 print(x_train.shape)
193 print "x_validate shape: ",
194 print(x_validate.shape)
195 print("")
196
197 # ===== #

```

Listing 3: Preprocessing the training data using the VGG-16 pretrained model and reshaping the data

```

180 # ===== #
181
182 # create the model
183 model = Sequential()
184
185 model.add(Dense(1024, activation='relu', input_shape=(25088,)))
186 model.add(Dropout(0.5))
187 model.add(Dense(512, activation='relu'))
188 model.add(Dropout(0.5))
189 model.add(Dense(256, activation='relu'))
190 model.add(Dropout(0.5))
191 model.add(Dense(256, activation='relu'))
192 model.add(Dropout(0.5))
193 model.add(Dense(128, activation='relu'))
194 model.add(Dropout(0.5))
195
196 model.add(Dense(3, activation='softmax'))
197
198 # ===== #
199
200 # create a checkpoint file to store the trained weights
201 mcp_save = ModelCheckpoint('weights.hdf5', save_best_only=True, monitor='val_loss', mode='min')
202
203 # compile the model
204 model.compile(loss='categorical_crossentropy', optimizer='Adam', metrics=['accuracy'])
205
206 # train the model
207 model.fit(x_train, y_train, epochs=25, validation_data=(x_validate, y_validate), callbacks=[mcp_save], batch_size=128)
208
209 # remove used frames in the extracted_frames folder
210 files = glob('training_videos/extracted_frames/*.jpg')
211 for f in files:
212     os.remove(f)

```

Listing 4: Creating, compiling and training the model

```

149 # ===== #
150
151 # load the frame list
152 test = pd.read_csv('testing_frames_list.csv')
153 test.head()
154
155 actual = test['class']
156 predictions = []
157
158 # extract video frames and make prediction
159 for i in tqdm(range(test.shape[0])):
160     # loading the image and keeping the target size as (224,224,3)
161     img = image.load_img('testing_videos/extracted_frames/'+test['image'][i], target_size=(224,224,3))
162     # converting it to array
163     img = image.img_to_array(img)
164     # normalizing the pixel value
165     img = img/255
166
167     # preprocess with VGG-16 base model and reshape
168     test_image = []
169     test_image.append(img)
170     x_test = np.array(test_image)
171     x_test = base_model.predict(x_test)
172     x_test = x_test.reshape(x_test.shape[0], 7*7*512)
173
174     # make prediction using our trained model
175     prediction = model.predict_classes(x_test) # 0 == housecleaning, 1 == not_housework, 2 == washingdishes
176     probability = model.predict_proba(x_test) # [1,0,0] == housecleaning, [0,1,0] == not_housework, [0,0,1] == washingdishes
177     predictions.append(1.0 - probability[0][1]) # probability[0][1] == probability of not_housework
178     if prediction == 0 and actual[i] == "housecleaning":
179         num_total_correct += 1
180         correctness_list.append('yes')
181     elif prediction == 1 and actual[i] == "not_housework":
182         num_total_correct += 1
183         correctness_list.append('yes')
184     elif prediction == 2 and actual[i] == "washingdishes":
185         num_total_correct += 1
186         correctness_list.append('yes')
187     else:
188         correctness_list.append('no')
189     num_frames += 1
190     predictions_list.append(prediction)
191     probability_list.append(1.0 - probability[0][1])
192
193 # combine time/label data and write out to JSON file
194 for i in range(len(timestamps)):
195     data_outfile.append([str(timestamps[i]), str(predictions[i])])
196 plt.ylim(-0.25,1.25)
197 plt.plot(timestamps, predictions)
198 plt.savefig(videoFile + '.png')
199 plt.close()
200
201 with open(videoFile + '.json', 'w') as outfile:
202     json.dump(data_outfile, outfile)

```

Listing 5: Extracting frames from the testing set and feeding to the model

7. Training and Testing Performance

```
phatnguyen@27-MAR-20: ~/Documents/CSCE-636-Video-Classification
File Edit View Search Terminal Help
4992/8932 [=====>.....] - ETA: 4s - loss: 0.0511 - accuracy:
5120/8932 [=====>.....] - ETA: 4s - loss: 0.0501 - accuracy:
5248/8932 [=====>.....] - ETA: 4s - loss: 0.0503 - accuracy:
5376/8932 [=====>.....] - ETA: 4s - loss: 0.0525 - accuracy:
5504/8932 [=====>.....] - ETA: 4s - loss: 0.0529 - accuracy:
5632/8932 [=====>.....] - ETA: 3s - loss: 0.0522 - accuracy:
5760/8932 [=====>.....] - ETA: 3s - loss: 0.0523 - accuracy:
5888/8932 [=====>.....] - ETA: 3s - loss: 0.0523 - accuracy:
6016/8932 [=====>.....] - ETA: 3s - loss: 0.0517 - accuracy:
6144/8932 [=====>.....] - ETA: 3s - loss: 0.0515 - accuracy:
6272/8932 [=====>.....] - ETA: 3s - loss: 0.0524 - accuracy:
6400/8932 [=====>.....] - ETA: 2s - loss: 0.0516 - accuracy:
6528/8932 [=====>.....] - ETA: 2s - loss: 0.0511 - accuracy:
6656/8932 [=====>.....] - ETA: 2s - loss: 0.0512 - accuracy:
6784/8932 [=====>.....] - ETA: 2s - loss: 0.0514 - accuracy:
6912/8932 [=====>.....] - ETA: 2s - loss: 0.0513 - accuracy:
7040/8932 [=====>.....] - ETA: 2s - loss: 0.0513 - accuracy:
7168/8932 [=====>.....] - ETA: 2s - loss: 0.0507 - accuracy:
7296/8932 [=====>.....] - ETA: 1s - loss: 0.0508 - accuracy:
7424/8932 [=====>.....] - ETA: 1s - loss: 0.0502 - accuracy:
7552/8932 [=====>.....] - ETA: 1s - loss: 0.0495 - accuracy:
7680/8932 [=====>.....] - ETA: 1s - loss: 0.0501 - accuracy:
7808/8932 [=====>.....] - ETA: 1s - loss: 0.0498 - accuracy:
7936/8932 [=====>.....] - ETA: 1s - loss: 0.0492 - accuracy:
8064/8932 [=====>.....] - ETA: 1s - loss: 0.0492 - accuracy:
8192/8932 [=====>.....] - ETA: 0s - loss: 0.0496 - accuracy:
8320/8932 [=====>.....] - ETA: 0s - loss: 0.0500 - accuracy:
8448/8932 [=====>.....] - ETA: 0s - loss: 0.0501 - accuracy:
8576/8932 [=====>.....] - ETA: 0s - loss: 0.0497 - accuracy:
8704/8932 [=====>.....] - ETA: 0s - loss: 0.0492 - accuracy:
8832/8932 [=====>.....] - ETA: 0s - loss: 0.0495 - accuracy:
8932/8932 [=====>.....] - 11s 1ms/step - loss: 0.0491 - accur
acv: 0.9865 - val loss: 0.0325 - val accuracy: 0.9919
```

Listing 6: Validating accuracy 99.19%

```
phatnguyen@27-MAR-20: ~/Documents/CSCE-636-Video-Classification
File Edit View Search Terminal Help
100%|██████████| 963/967 [09:29<00:03, 1.17it/s]

Video File: WipingWindows-0183C.mp4
100%|██████████| 7/7 [00:00<00:00, 206761.46it/s]
100%|██████████| 7/7 [00:00<00:00, 18.22it/s]
100%|██████████| 964/967 [09:30<00:02, 1.17it/s]

Video File: WipingWindows-0184C.mp4
100%|██████████| 7/7 [00:00<00:00, 233016.89it/s]
100%|██████████| 7/7 [00:00<00:00, 18.37it/s]
100%|██████████| 965/967 [09:31<00:01, 1.17it/s]

Video File: WipingWindows-0185C.mp4
100%|██████████| 7/7 [00:00<00:00, 225847.14it/s]
100%|██████████| 7/7 [00:00<00:00, 18.83it/s]
100%|██████████| 966/967 [09:32<00:00, 1.16it/s]

Video File: WipingWindows-0186C.mp4
100%|██████████| 7/7 [00:00<00:00, 225847.14it/s]
100%|██████████| 7/7 [00:00<00:00, 18.24it/s]
100%|██████████| 967/967 [09:32<00:00, 1.69it/s]
('Total num frames: ', 7110)
('Correctly predicted:', 6339)
('Test Accuracy: ', 89.15611814345992)
phatnguyen@27-MAR-20:~/Documents/CSCE-636-Video-Classification$
```

Listing 7: Testing accuracy 89.16%

8. Improvements over Last Submission

```
phatnguyen@27-MAR-20: ~/Documents/CSCE-636-Video-Classification
File Edit View Search Terminal Help
phatnguyen@27-MAR-20:~/Documents/CSCE-636-Video-Classification$ python testing-s
ub6.py
Using TensorFlow backend.
WARNING:tensorflow:From /usr/local/lib/python2.7/dist-packages/tensorflow/python
/ops/nn_impl.py:180: where (from tensorflow.python.ops.array_ops) is deprecated
and will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
0%|          | 0/1 [00:00<?, ?it/s]

Video File: Combine.mp4
100%|          | 141/141 [00:00<00:00, 838860.80it/s]
100%|          | 141/141 [00:06<00:00, 22.09it/s]
100%|          | 1/1 [00:11<00:00, 11.60s/it]t/s]
('Total num frames: ', 141)
('Correctly predicted:', 0)
('Test Accuracy: ', 0.0)
phatnguyen@27-MAR-20:~/Documents/CSCE-636-Video-Classification$
```

Submission 6 (0%)

```
phatnguyen@27-MAR-20: ~/Documents/CSCE-636-Video-Classification
File Edit View Search Terminal Help
phatnguyen@27-MAR-20:~/Documents/CSCE-636-Video-Classification$ python testing.p
y
Using TensorFlow backend.
0%|          | 0/1 [00:00<?, ?it/s]

Video File: Combine.mp4
100%|          | 141/141 [00:00<00:00, 783307.10it/s]
100%|          | 141/141 [00:09<00:00, 14.85it/s]
100%|          | 1/1 [00:15<00:00, 15.44s/it]t/s]
('Total num frames: ', 141)
('Correctly predicted:', 141)
('Test Accuracy: ', 100.0)
phatnguyen@27-MAR-20:~/Documents/CSCE-636-Video-Classification$
```

Submission 8 (100%)

Listing 8: Accuracy on YouTube sample video provided by professor Jiang

```
phatnguyen@27-MAR-20: ~/Documents/CSCE-636-Video-Classification
File Edit View Search Terminal Help
Video File: YoYo_g05_c02.avi
100%|          | 8/8 [00:00<00:00, 138084.08it/s]
100%|          | 8/8 [00:00<00:00, 21.12it/s]
99%|          | 499/502 [05:00<00:01, 2.25it/s]

Video File: YoYo_g05_c03.avi
100%|          | 8/8 [00:00<00:00, 170327.07it/s]
100%|          | 8/8 [00:00<00:00, 21.53it/s]
100%|          | 500/502 [05:01<00:00, 2.24it/s]

Video File: YoYo_g05_c04.avi
100%|          | 7/7 [00:00<00:00, 199728.76it/s]
100%|          | 7/7 [00:00<00:00, 21.51it/s]
100%|          | 501/502 [05:01<00:00, 2.32it/s]

Video File: YoYo_g05_c05.avi
100%|          | 8/8 [00:00<00:00, 241398.79it/s]
100%|          | 8/8 [00:00<00:00, 21.59it/s]
100%|          | 502/502 [05:01<00:00, 1.66it/s]
('Total num frames: ', 4064)
('Correctly predicted:', 2050)
('Test Accuracy: ', 50.44291338582677)
phatnguyen@27-MAR-20:~/Documents/CSCE-636-Video-Classification$
```

Submission 6 (50.44%)

```
phatnguyen@27-MAR-20: ~/Documents/CSCE-636-Video-Classification
File Edit View Search Terminal Help
Video File: YoYo_g05_c01.avi
100%|          | 8/8 [00:00<00:00, 250406.21it/s]
100%|          | 8/8 [00:00<00:00, 19.52it/s]
99%|          | 498/502 [05:15<00:01, 2.12it/s]

Video File: YoYo_g05_c02.avi
100%|          | 8/8 [00:00<00:00, 258111.02it/s]
100%|          | 8/8 [00:00<00:00, 19.40it/s]
99%|          | 499/502 [05:15<00:01, 2.11it/s]

Video File: YoYo_g05_c03.avi
100%|          | 8/8 [00:00<00:00, 250406.21it/s]
100%|          | 8/8 [00:00<00:00, 19.17it/s]
100%|          | 500/502 [05:16<00:00, 2.10it/s]

Video File: YoYo_g05_c04.avi
100%|          | 7/7 [00:00<00:00, 233016.89it/s]
100%|          | 7/7 [00:00<00:00, 19.14it/s]
100%|          | 501/502 [05:16<00:00, 2.16it/s]

Video File: YoYo_g05_c05.avi
100%|          | 8/8 [00:00<00:00, 250406.21it/s]
100%|          | 8/8 [00:00<00:00, 19.12it/s]
100%|          | 502/502 [05:17<00:00, 1.58it/s]
('Total num frames: ', 4109)
('Correctly predicted:', 3318)
('Test Accuracy: ', 80.7495741056218)
phatnguyen@27-MAR-20:~/Documents/CSCE-636-Video-Classification$
```

Submission 8 (80.75%)

Listing 9: Accuracy on 19 unmet/untrained activities

9. Instructions on How to Test the Trained DNN

- Dependencies:
 - Python 2.7
 - Keras
 - Tensorflow
 - OpenCV
 - Scipy, sklearn, skimage, glob, tqdm
- How to train:
 - Put the full name of the training videos in the *trainlist.txt* file
 - Put the training video files in the *training_videos* folder
 - Start the training process using command: `python training.py`
- How to test:
 - Put the full name of the training videos in the *testlist.txt* file
 - Put the testing video files in the *testing_videos* folder
 - Run the test using command: `python testing.py`
- For testing script and other instructions, please visit the GitHub page of this project:
<https://github.com/phatnguyen0430/CSCE-636-Video-Classification>

10. Reference

[1] Step-by-Step Deep Learning Tutorial to Build your own Video Classification Model.
<https://www.analyticsvidhya.com/blog/2019/09/step-by-step-deep-learning-tutorial-video-classification-python/>